

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

Customers = pd.read_csv("/content/Customers.csv")

Subscriptions = pd.read_csv("/content/Subscription.csv")

Transactions = pd.read_csv("/content/Transcation.csv")

Usage = pd.read_csv("/content/Usage.csv")

merged_df = pd.merge(Customers, Subscriptions, on="CustomerID", how="left")

merged_df = merged_df.merge(Transactions, on="CustomerID", how="left")

merged_df = merged_df.merge(Usage, on="CustomerID", how="left")

# Check for non-datetime values
print(merged_df[["StartDate", "EndDate"]].dtypes)

# Look for rows with invalid dates
print(merged_df[merged_df["StartDate"].isna() | merged_df["EndDate"].isna()])
```



68	15.99	Credit	5.0	Watch Movie	2024-02-08
69	15.99	Credit	30.0	Rate Content	2024-02-25

	usage_amount
0	NaN
2	21.23
7	30.51
9	55.69
11	76.66
12	76.66
16	NaN
18	78.63
19	78.63
21	54.42
22	54.42
24	55.69
25	55.69
32	87.98
33	50.84
35	NaN
38	33.81
41	36.39
42	67.21
45	NaN
46	NaN
47	NaN
49	NaN
51	NaN
55	NaN
59	47.97
60	50.43
61	51.60
62	40.62
63	63.91
65	68.99
66	75.87
68	27.90
69	32.63

```
print(merged_df.columns)
```

```
Index(['CustomerID', 'Name', 'Age', 'Gender', 'Income', 'Location',
      'SubscriptionID', 'StartDate', 'EndDate', 'Status', 'transaction_id',
      'transaction_date', 'amount', 'transaction_type', 'usage_id',
      'feature_used', 'usage_date', 'usage_amount'],
      dtype='object')
```

```
merged_df["StartDate"] = pd.to_datetime(merged_df["StartDate"], errors="coerce")
merged_df["EndDate"] = pd.to_datetime(merged_df["EndDate"], errors="coerce")
merged_df["tenure"] = (merged_df["EndDate"] - merged_df["StartDate"]).dt.days
```

```
print(merged_df[["StartDate", "EndDate", "tenure"]].head())
```

```

⇨
   StartDate  EndDate  tenure
0 2022-01-31      NaT      NaN
1 2022-02-28 2024-11-21  997.0
2 2022-03-31      NaT      NaN
3 2022-04-30 2024-11-21  936.0
4 2022-04-30 2024-11-21  936.0

```

```

# Avoid division by zero by replacing tenure of 0 with NaN
merged_df["tenure"] = merged_df["tenure"].replace(0, np.nan)

```

```

# Calculate average monthly spend
merged_df["average_monthly_spend"] = merged_df["amount"] / (merged_df["tenure"] / 30.0)

```

```

# Fill NaN values in 'average_monthly_spend' if any tenure is missing or invalid
merged_df["average_monthly_spend"].fillna(0, inplace=True)

```

```

⇨ <ipython-input-78-fca80104520c>:8: FutureWarning: A value is trying to be set on a copy
The behavior will change in pandas 3.0. This inplace method will never work because the

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col

```

```
merged_df["average_monthly_spend"].fillna(0, inplace=True)
```

```
print(merged_df.columns)
```

```

⇨ Index(['CustomerID', 'Name', 'Age', 'Gender', 'Income', 'Location',
        'SubscriptionID', 'StartDate', 'EndDate', 'Status', 'transaction_id',
        'transaction_date', 'amount', 'transaction_type', 'usage_id',
        'feature_used', 'usage_date', 'usage_amount', 'tenure',
        'average_monthly_spend'],
        dtype='object')

```

```
print(merged_df.head())
```

```

⇨
   CustomerID  Name  Age  Gender  Income  Location  SubscriptionID \
0           1  Karan Reddy  29   Male   40000  Hyderabad      SUB0001
1           2  Anaya Joshi  35  Female   55000  Bangalore      SUB0002
2           3    Diya Das  28  Female   60000   Chennai      SUB0003
3           4 Siddharth Reddy  32   Male   45000    Delhi      SUB0004
4           4 Siddharth Reddy  32   Male   45000    Delhi      SUB0004

   StartDate  EndDate  Status  transaction_id  transaction_date  amount \
0 2022-01-31      NaT    Yes              NaN              NaN      NaN
1 2022-02-28 2024-11-21    No      e183937c      2024-10-01    13.92
2 2022-03-31      NaT    Yes              NaN              NaN      NaN


```

3	2022-04-30	2024-11-21	No	ce657026	2024-06-05	17.95
4	2022-04-30	2024-11-21	No	ce657026	2024-06-05	17.95

	transaction_type	usage_id	feature_used	usage_date	usage_amount	\
0	NaN	NaN	NaN	NaN	NaN	
1	Debit	4.0	Watch Movie	2024-08-02	30.94	
2	NaN	11.0	Rate Content	2024-02-21	21.23	
3	Debit	21.0	Download Episode	2024-11-05	82.45	
4	Debit	28.0	Watch Movie	2024-09-04	66.47	

	tenure	average_monthly_spend
0	NaN	0.000000
1	997.0	0.418857
2	NaN	0.000000
3	936.0	0.575321
4	936.0	0.575321

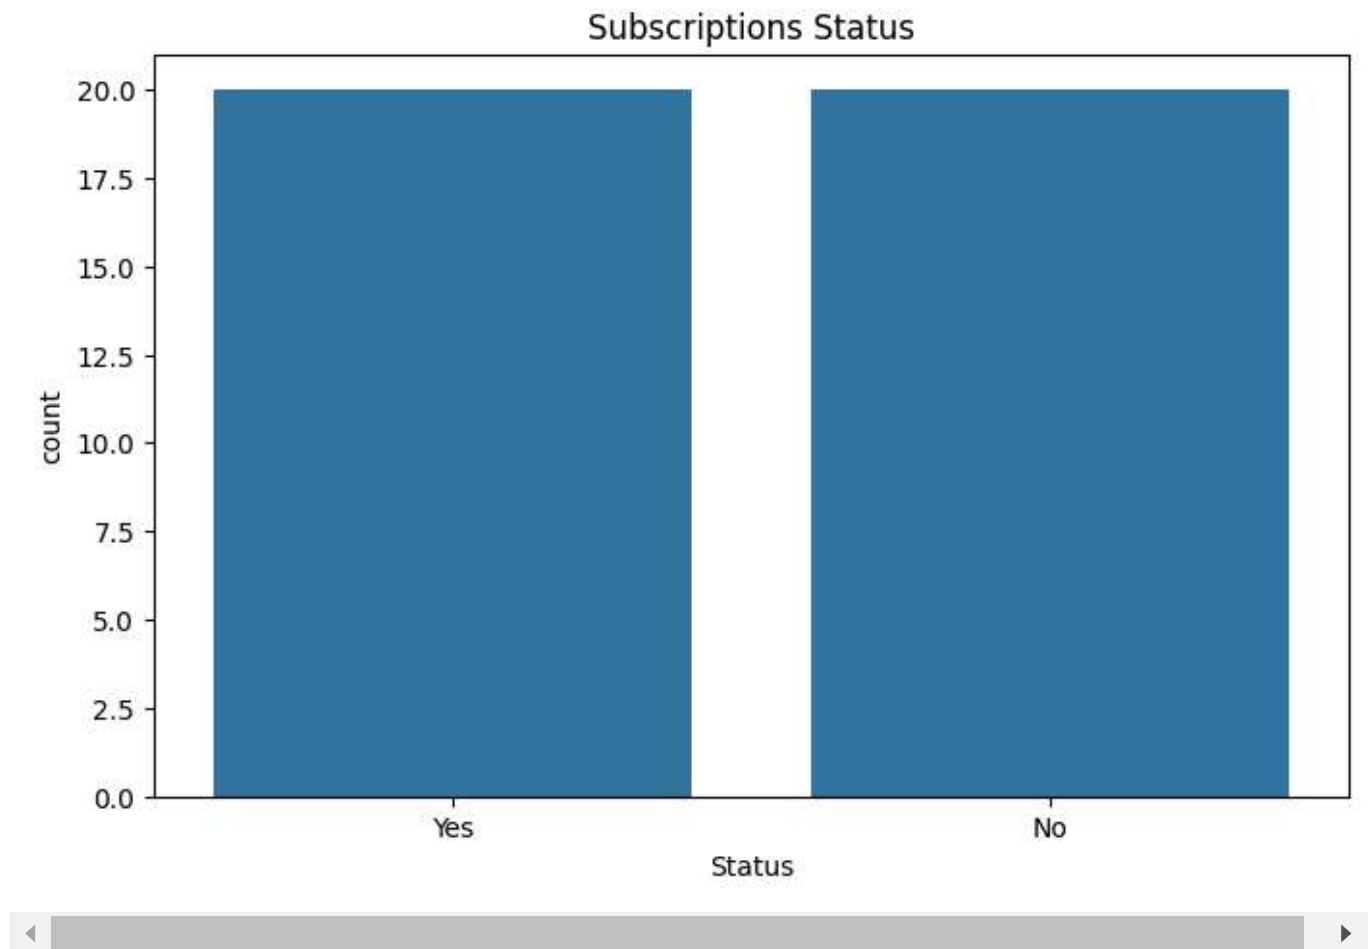
```
plt.figure(figsize=(8, 5))
sns.histplot(merged_df["Age"], bins=10, kde=True)
plt.title("Age Distribution")
plt.show()
```



```
Subscriptions = pd.read_csv("/content/Subscription.csv")
print(Subscriptions.columns)
plt.figure(figsize=(8, 5))
```

```
sns.countplot(x="Status", data=Subscriptions)  
plt.title("Subscriptions Status")  
plt.show()
```


➞ Index(['CustomerID', 'SubscriptionID', 'StartDate', 'EndDate', 'Status'], dtype='object')



```
print(Subscriptions.columns)
```

➞ Index(['CustomerID', 'SubscriptionID', 'StartDate', 'EndDate', 'Status'], dtype='object')

```
plt.figure(figsize=(8, 5))  
sns.countplot(x="transaction_type", data=Transactions)  
plt.title("Transaction Types")  
plt.show()
```



```
import pandas as pd
from sklearn.preprocessing import LabelEncoder

# Print columns to inspect available ones
print(merged_df.columns)

# Apply label encoding with correct column names
merged_df["Gender"] = LabelEncoder().fit_transform(merged_df["Gender"])
merged_df["Location"] = LabelEncoder().fit_transform(merged_df["Location"])

# Use the correct column names based on your DataFrame
merged_df["SubscriptionID"] = LabelEncoder().fit_transform(merged_df["SubscriptionID"]) # 1
merged_df["transaction_type"] = LabelEncoder().fit_transform(merged_df["transaction_type"])
```



```
Index(['CustomerID', 'Name', 'Age', 'Gender', 'Income', 'Location',
      'SubscriptionID', 'StartDate', 'EndDate', 'Status', 'transaction_id',
      'transaction_date', 'amount', 'transaction_type', 'usage_id',
      'feature_used', 'usage_date', 'usage_amount', 'tenure',
      'average_monthly_spend'],
      dtype='object')
```

Start coding or [generate](#) with AI.

```

# Check column names to make sure 'churned' is present
print(merged_df.columns)

# If 'churned' column doesn't exist, create it
if 'churned' not in merged_df.columns:
    merged_df["churned"] = (merged_df["Status"] == "Churned").astype(int)

# Define features and target variable
features = ["Age", "Gender", "Income", "tenure", "average_monthly_spend"]
X = merged_df[features]
y = merged_df["churned"]

Index(['CustomerID', 'Name', 'Age', 'Gender', 'Income', 'Location',
      'SubscriptionID', 'StartDate', 'EndDate', 'Status', 'transaction_id',
      'transaction_date', 'amount', 'transaction_type', 'usage_id',
      'feature_used', 'usage_date', 'usage_amount', 'tenure',
      'average_monthly_spend'],
      dtype='object')

from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report

# Split data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Train model
model = RandomForestClassifier(random_state=42)
model.fit(X_train, y_train)

# Evaluate model
predictions = model.predict(X_test)
print(classification_report(y_test, predictions))

```

```

precision    recall  f1-score   support

0           1.00      1.00      1.00         15

accuracy               1.00         15
macro avg              1.00      1.00      1.00         15
weighted avg           1.00      1.00      1.00         15

```

```

from sklearn.impute import SimpleImputer

# Create an imputer to fill missing values with the mean
imputer = SimpleImputer(strategy='mean')

# Impute the missing values in X

```

```
X_imputed = imputer.fit_transform(X)

# Perform clustering on the cleaned data
kmeans = KMeans(n_clusters=3, random_state=42)
merged_df["segment"] = kmeans.fit_predict(X_imputed)

# Analyze clusters
for segment in merged_df["segment"].unique():
    print(merged_df[merged_df["segment"] == segment].describe())
```



50%	15.750000	0.000000	22.000000	54.420000	326.000000
75%	16.307500	1.000000	30.250000	76.660000	326.000000
max	18.900000	2.000000	36.000000	78.630000	326.000000
std	3.941125	0.785905	12.802403	21.668911	211.310199

	average_monthly_spend	churned	segment
count	17.000000	17.0	17.0
mean	0.081577	0.0	1.0
min	0.000000	0.0	1.0
25%	0.000000	0.0	1.0
50%	0.000000	0.0	1.0
75%	0.000000	0.0	1.0
max	0.783129	0.0	1.0
std	0.232458	0.0	0.0

```
Transactions["transaction_date"] = pd.to_datetime(Transactions["transaction_date"])
```

```
# Prepare revenue data
```

```
revenue_data = Transactions.groupby(Transactions["transaction_date"].dt.to_period("M"))["amc"]
revenue_data.columns = ["ds", "y"]
revenue_data["ds"] = revenue_data["ds"].dt.to_timestamp()
```

```
# Fit Prophet model
```

```
model = Prophet()
model.fit(revenue_data)
```

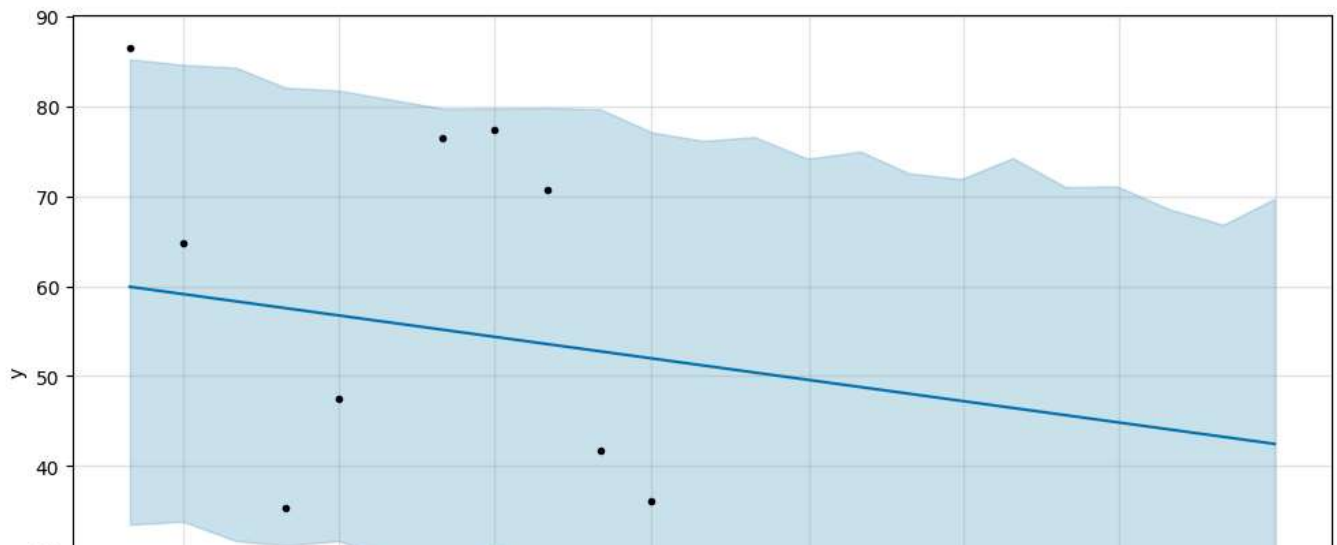
```
# Forecast revenue
```

```
future = model.make_future_dataframe(periods=12, freq="M")
forecast = model.predict(future)
model.plot(forecast)
plt.show()
```

```

INFO:prophet:Disabling yearly seasonality. Run prophet with yearly_seasonality=True to c
INFO:prophet:Disabling weekly seasonality. Run prophet with weekly_seasonality=True to c
INFO:prophet:Disabling daily seasonality. Run prophet with daily_seasonality=True to ove
INFO:prophet:n_changepoints greater than number of observations. Using 7.
DEBUG:cmdstanpy:input tempfile: /tmp/tmpqyozztqq/uraxaakh.json
DEBUG:cmdstanpy:input tempfile: /tmp/tmpqyozztqq/o1wt7bpx.json
DEBUG:cmdstanpy:idx 0
DEBUG:cmdstanpy:running CmdStan, num_threads: None
DEBUG:cmdstanpy:CmdStan args: ['/usr/local/lib/python3.10/dist-packages/prophet/stan_moc
12:05:47 - cmdstanpy - INFO - Chain [1] start processing
INFO:cmdstanpy:Chain [1] start processing
12:05:47 - cmdstanpy - INFO - Chain [1] done processing
INFO:cmdstanpy:Chain [1] done processing
/usr/local/lib/python3.10/dist-packages/prophet/forecaster.py:1854: FutureWarning: 'M' i
dates = pd.date_range(

```



```

merged_df.to_csv("customer_analysis_results.csv", index=False)
forecast.to_csv("revenue_forecast.csv", index=False)

```

Start coding or [generate](#) with AI.